

MP6: Primitive Disk Device Driver

Cheng-Yun Cheng
UIN: 633002216
CSCE611: Operating System

Assigned Tasks

Main: Completed.

Bonus Option 1: Completed.

Bonus Option 2: Completed.

Bonus Option 3: Completed.

Bonus Option 4: Completed.

System Design

The goal of the machine problem is to implement a kernel-level device driver.

- Main part: When the read and write functions are called, the current thread would be put into a block queue and give up the CPU. Every time a thread yield the CPU, the scheduler would check the status of the disk. If the disk is ready, the scheduler will resume the first thread in the block queue (put it into the ready queue).
- Option 1, support for disk mirroring: A MirroredDisk class is implemented. When issuing a read operation, the data transfer would be done when one of the disks is ready. Write operations would be issued to both disks.
- Option 2, using interrupts for concurrency: Rather than checking the state of the disk at regular intervals, interrupt is used to put the thread which is waiting for transferring data into the ready queue. I found that the interrupt will occur after the write operation is finished. Therefore, I only used interrupt for read operation, and busy-waiting is used for write operation.
- Option 3, design of a thread-safe disk system: If multiple threads can access the disk, we should consider when to issue the operation. If there is already a thread issuing an I/O operation, then another thread cannot issue an operation at this time. It should issue the operation after the last one is finished. To do so, when a thread wants to execute an I/O operation, put it into the block queue and yield the CPU. A boolean variable `need_issue` is set to determine that whether the first thread in the block queue is able to call the `issue_operation` function. If `need_issue` is true, then dispatch to this thread. On the other hand, dispatch to the next thread in the ready queue. There is a simple example.
 1. Thread 2 issues an I/O operation, `need_issue` is set as false.
 2. Thread 2 waits for the disk and yields the CPU.
 3. Thread 3 wants to request I/O operation, but there is already an issue. Therefore, thread 3 yields.
 4. Disk is ready and Thread 2 transfers data. `need_issue` is set as true.
 5. Now Thread 3 can issue the operation.

Code Description (Main)

For the main part, I modified the `blocking_disk.h`, `blocking_disk.c`, and `kernel.c`. Also, `queue.h`, `scheduler.h`, and `scheduler.c` are added. Also, the `issue_operation` function in `SimpleDisk` is moved to protected part and modified as a virtual function.

kernel.c :

- `_USES_SCHEDULER_` is used.
- Include `blocking_disk.h`.
- `SimpleDisk` is replaced with `BlockingDisk`.

blocking_disk.h: Data member :

- Queue `block_queue`: A queue to store threads which are waiting for the disk.
- `wait_CPU`: If the disk is ready and the thread which issued I/O operation is waiting for execution in the ready queue, `wait_CPU` is set as true.

blocking_disk.c: `wait_until_ready` : Put the current thread into the block queue and call the yield function to give up the CPU.

```
void BlockingDisk::wait_until_ready(){
    block_queue.enqueue(Thread::CurrentThread());
    Console::puts("Wait for disk, yield\n");
    SYSTEM_SCHEDULER->yield();
}
```

blocking_disk.c: `disk_ready` : Return true if the block queue is not empty and the disk is ready.

```
bool BlockingDisk::disk_ready(){
    return SimpleDisk::is_ready() && !block_queue.is_empty();
}
```

blocking_disk.c: `is_wait_CPU` : Return `wait_CPU`.

```
bool BlockingDisk::is_wait_CPU(){
    return wait_CPU;
}
```

blocking_disk.c: `dequeue` : Set `wait_CPU` as true. Then pop the first thread in the block queue and return it.

```
Thread * BlockingDisk::dequeue(){
    wait_CPU = true;
    return block_queue.dequeue();
}
```

blocking_disk.c: `read` : Call the `read` function in `SimpleDisk`. After finishing it, set `wait_CPU` as false.

```
void BlockingDisk::read(unsigned long _block_no, unsigned char * _buf) {
    // -- REPLACE THIS!!!
    SimpleDisk::read(_block_no, _buf);
    wait_CPU = false;
    Console::puts("READ DONE\n");
}
```

blocking_disk.c: write : Call the write function in SimpleDisk. After finishing it, set wait_CPU as false.

```
void BlockingDisk::write(unsigned long _block_no, unsigned char * _buf) {
    // -- REPLACE THIS!!!
    SimpleDisk::write(_block_no, _buf);
    wait_CPU = false;
    Console::puts("WRITE DONE\n");
}
```

scheduler.h: Scheduler : The most part of Scheduler is as same as the Scheduler in MP5, except yield function.

- Data member: Queue ready_queue
- Member function: yield, resume, add, terminate

scheduler.c: yield : Check whether the disk is ready or not first. If yes, pop the thread from the block queue and put it into the ready queue. Then, get the next thread from the ready queue and call the dispatch_to function to invoke the context switch.

```
void Scheduler::yield() {
    if(SYSTEM_DISK->disk_ready() && !SYSTEM_DISK->is_wait_CPU()){
        resume(SYSTEM_DISK->dequeue());
        Console::puts("Disk is ready, put thread into ready queue\n");
    }

    Thread * next = ready_queue.dequeue();
    Thread::dispatch_to(next);
}
```

queue.h: Queue : Queue is implement by array. It is as same as the Queue class in MP5. The member functions include is_empty, enqueue, dequeue, and remove.

Testing (Main)

For the machine problem, I only use the provided test. Only thread 2 requests I/O operation. The result shows that after issuing the read and write operation, thread 2 yields the CPU. After the disk is ready, thread 2 is put into ready thread again to wait for transfer data.

```
FUN 1: TICK [8]
FUN 1: TICK [9]
THREAD: 1
FUN 2 INVOKED!
FUN 2 IN ITERATION[0]
Thread2 Reading a block from disk...
Wait for disk, yield
Disk is ready, put thread into ready queue
THREAD: 2
FUN 3 INVOKED!
FUN 3 IN BURST[0]
FUN 3: TICK [0]
```

[illegible]

Code Description (Option 1)

I added the `MirroredDisk` class in `blocking_disk.h` and `blocking_disk.c`.

blocking_disk.h: MirroredDisk : The MirroredDisk class is derived from the SimpleDisk class, and most functions in MirroredDisk is as same as the functions in BlockingDisk.

```
class MirroredDisk : public SimpleDisk{
private:
    Queue block_queue;
    bool wait_CPU;
protected:
    virtual void issue_operation(DISK_OPERATION_op, unsigned long_block_no, DISK_ID_disk_id);
    virtual void wait_until_ready();
    virtual void write(unsigned long_block_no, unsigned char * _buf, DISK_ID_disk_id);
public:
    MirroredDisk(unsigned int _size);
    bool disk_ready();
    bool is_wait_CPU();
    Thread * dequeue();
    virtual void read(unsigned long_block_no, unsigned char * _buf);
    virtual void write(unsigned long_block_no, unsigned char *_buf);
};
```

blocking_disk.c: issue_operation : The function is as same as the issue_operation function in SimpleDisk, except a DISK_ID argument. This argument is used to determine which disk we will issue operation to.

```
void MirrorDisk::issueOperation(DISK_OPERATION op, unsigned long block_no, DISK_ID disk_id) {
    Machine::outportb(0x1F, 0x80); /* send NULL to port 0x1F */
    Machine::outportb(0x1E, 0x01); /* send sector count to port 0x1E */
    Machine::outportb(0x1F, (unsigned char)block_no); /* send char block number */
    Machine::outportb(0x1F, (unsigned char)(block_no >> 8)); /* send low 8 bits of block number */
    Machine::outportb(0x1F, (unsigned char)(block_no >> 16)); /* send next 8 bits of block number */
    Machine::outportb(0x1F, (unsigned char)(block_no >> 24)); /* send next 8 bits of block number */
    unsigned int disk_no = disk_id == DISK_ID_MASTER ? 0 : 1;
    Machine::outportb(0x1F, ((unsigned char)(block_no >> 24)&0xF) | 0xE0 | (disk_no << 4));
    Machine::outportb(0x1F, (/* send drive indicator, some bits, highest 4 bits of block no */
        (op == DISK_OPERATION::READ) ? 0x20 : 0x30));
}
```

blocking_disk.c: _write : A DISK_ID argument is added to determine which disk would do the operation..

```
void MirroredDisk::write(unsigned long block no, unsigned char * _buf, DISK_ID _disk id)
{
    issue_operation(DISK_OPERATION::WRITE, _block no, _disk id);
    wait_until_ready();
    int i;
    unsigned short tmpwp;
    for (i = 0; i < 256; i++) {
        tmpwp = _buf[2*i] | (_buf[2*i+1] << 8);
        Machine::outportw(0x1f0, tmpwp);
    }
}
```

blocking_disk.c: read : The difference between this function and the read function in SimpleDisk is that the read operation is issued to both disks.

```
void MirroredDisk::read(unsigned long _block_no, unsigned char * _buf){
    issue_operation(DISK_OPERATION::READ, _block_no, DISK_ID::MASTER);
    issue_operation(DISK_OPERATION::READ, _block_no, DISK_ID::DEPENDENT);

    wait_until_ready();

    int i;
    unsigned short tmpw;
    for (i = 0; i < 256; i++) {
        tmpw = Machine::inportw(0x1F0);
        _buf[i*2] = (unsigned char)tmpw;
        _buf[i*2+1] = (unsigned char)(tmpw >> 8);
    }
    wait_CPU = false;
    Console::puts("READ DONE\n");
}
```

blocking_disk.c: write : Call _write function for the master disk and the dependent disk.

```
void MirroredDisk::write(unsigned long _block_no, unsigned char * _buf){
    _write( _block_no, _buf, DISK_ID::MASTER);
    wait_CPU = false;
    Console::puts("WRITE MASTER DONE\n");

    _write( _block_no, _buf, DISK_ID::DEPENDENT);
    wait_CPU = false;
    Console::puts("WRITE DEPENDENT DONE\n");
}
```

Testing (Option 1)

I only use the provided test. To do this test, please uncomment `_MIRRORED_DISK` in `kernel.c` and `scheduler.c`.

```
FUN 1: TICK [9]
THREAD: 1
FUN 2 INVOKED!
FUN 2 IN ITERATION[0]
Thread2 Reading a block from disk...
Wait for disk, yield
Disk is ready, put thread into ready queue
THREAD: 2
FUN 3 INVOKED!
```

[illegible]

```
FUN 1: TICK [8]
FUN 1: TICK [9]
WRITE MASTER DONE
Wait for disk, yield
Disk is ready, put thread into ready queue
FUN 3 IN BURST[2]
FUN 3: TICK [0]
```

```

FUN 1: TICK [8]
FUN 1: TICK [9]
WRITE DEPENDENT DONE
Thread2 finish
FUN 3 IN BURST[3]
FUN 3: TICK [0]

```

Code Description (Option 2)

I changed scheduler.c, blocking_disk.h, and blocking_disk.c. Also, an interrupt handler is added in kernel.c and enable_interrupt is added in thread.c.

blocking_disk.h: data member : bool done is added. It is used to determine that when interrupt 14 occurs, the data transfer is done or not. For read operation, interrupt occurs before transferring data. However, for write operation, interrupt occurs after transferring data.

blocking_disk.c: is_done : Return done.

```

bool BlockingDisk::is_done(){
    return done;
}

```

blocking_disk.c: read : Most parts are the same. In the read function, interrupt is disabled. done is set as false after issue_operation and set as true after data transfer.

```

void BlockingDisk::read(unsigned long _block_no, unsigned char * _buf) {
#ifdef INTERRUPT
    bool enable = false;
    if(Machine::interrupts_enabled()){
        Machine::disable_interrupts();
        enable = true;
    }
#endif

    SimpleDisk::issue_operation(DISK_OPERATION::READ, _block_no);
    done = false;

    wait_until_ready();

    /* read data from port */
    int i;
    unsigned short tmpw;
    for (i = 0; i < 256; i++) {
        tmpw = Machine::inportw(0x1F0);
        _buf[i*2] = (unsigned char)tmpw;
        _buf[i*2+1] = (unsigned char)(tmpw >> 8);
    }

    wait_CPU = false;
    done = true;
    Console::puts("READ DONE\n");

#ifdef INTERRUPT
    if(enable)
        Machine::enable_interrupts();
#endif
}

```

blocking_disk.c: write : Most parts are the same. In the write function, interrupt is disabled. done is set as false after issue_operation and set as true after data transfer. Besides, busy_waiting is used.

```

void BlockingDisk::write(unsigned long _block_no, unsigned char * _buf) {
#ifdef _INTERRUPT
    bool enable = false;
    if(Machine::interrupts_enabled()){
        Machine::disable_interrupts();
        enable = true;
    }
#endif

    SimpleDisk::issue_operation(DISK_OPERATION::WRITE, _block_no);
    done = false;

    while(!is_ready()){}

    /* write data to port */
    int i;
    unsigned short tmpw;
    for (i = 0; i < 256; i++) {
        tmpw = _buf[2*i] | (_buf[2*i+1] << 8);
        Machine::outportw(0x1F0, tmpw);
    }

    wait_CPU = false;
    done = true;
    Console::puts("WRITE DONE\n");

#ifdef _INTERRUPT
    if(enable)
        Machine::enable_interrupts();
#endif
}

```

scheduler.c: yield : Interrupt is disabled, and just pop a thread in the ready queue and dispatch to it.

```

void Scheduler::yield() {
#ifdef _INTERRUPT
    bool enable = false;
    if(Machine::interrupts_enabled()){
        Machine::disable_interrupts();
        enable = true;
    }
#endif

    Thread * next = ready_queue.dequeue();
    Thread::dispatch_to(next);

#ifdef _INTERRUPT
    if(enable)
        Machine::enable_interrupts();
#endif
}

```

kernel.c: DiskHandler : It is used for handling interrupt 14. When interrupt occurs, if done is false (read operation), pop a thread in the block queue and add it into the ready queue. If done is true (write operation), do nothing.

```

#ifdef _INTERRUPT
/* interrupt handler for disk */
class DiskHandler : public InterruptHandler{
public:
    virtual void handle_interrupt(REGS * _r) {
        Console::puts("Disk Interrupt Handler\n");
        if(!((BlockingDisk*)SYSTEM_DISK)->is_done()){
            Console::puts("put into ready queue\n");
            Thread * t = ((BlockingDisk*)SYSTEM_DISK)->dequeue();
            SYSTEM_SCHEDULER->resume(t);
        }
    }
} disk_handler;
InterruptHandler::register_handler(14, &disk_handler);
#endif

```

Testing (Option 2)

I only use the provided test. To do this test, please uncomment the `_INTERRUPT` in `kernel.c`, `scheduler.c`, `blocking_disk.c`, and `thread.c`.

```
FUN 1: TICK [8]
FUN 1: TICK [9]
THREAD: 1
FUN 2 INVOKED!
FUN 2 IN ITERATION[0]
Thread2 Reading a block from disk...
Wait for disk, yield
Disk Interrupt Handler
put into ready queue
THREAD: 2
FUN 3 INVOKED!
FUN 3 IN BURST[0]
FUN 3: TICK [0]
```

[illegible]

Code Description (Option 4)

I changed scheduler.c, blocking_disk.h, and blocking_disk.c.

blocking_disk.h: data member : bool issued is added. It means that whether an operation is issued.

blocking_disk.c: need issued : Return true, if there is not an operation issue and there is a thread in the block queue. It means that the the first thread in the block queue needs to issue an operation.

```
bool BlockingDisk::need_issued(){
    return !issued && !block_queue.is_empty();
}
```

blocking_disk.c: get_head : Get and return the first thread in the block queue (not pop, it is still in the queue).


```
Thread * BlockingDisk::get_head(){
    return block_queue.get_head();
}
```

blocking_disk.c: read :

- First, add the current thread into the block queue and yield the CPU.
- Get CPU again (when it is time for issue operation), issue the operation and set issued as true. Then yield the CPU.
- Get CPU again (when the disk is ready), transfer the data. After that, set issued as false.

```
void BlockingDisk::read(unsigned long _block_no, unsigned char * _buf) {
    block_queue.enqueue(Thread::CurrentThread());
    Console::puts("Wait for issuing operation, yield\n");
    SYSTEM_SCHEDULER->yield();

    SimpleDisk::issue_operation(DISK_OPERATION::READ, _block_no);
    issued = true;
    done = false;

    Console::puts("Wait for disk, yield\n");
    SYSTEM_SCHEDULER->yield();

    /* read data from port */
    int i;
    unsigned short tmpw;
    for (i = 0; i < 256; i++) {
        tmpw = Machine::inportw(0x1F0);
        _buf[i*2] = (unsigned char)tmpw;
        _buf[i*2+1] = (unsigned char)(tmpw >> 8);
    }

    wait_CPU = false;
    done = true;
    issued = false;
    Console::puts("READ DONE\n");
}
```

blocking_disk.c: write : Like read function.

- First, add the current thread into the block queue and yield the CPU.
- Get CPU again (when it is time for issue operation), issue the operation and set issued as true. Then yield the CPU.
- Get CPU again (when the disk is ready), transfer the data. After that, set issued as false.

```
void BlockingDisk::write(unsigned long _block_no, unsigned char * _buf) {
    block_queue.enqueue(Thread::CurrentThread());
    Console::puts("Wait for issuing operation, yield\n");
    SYSTEM_SCHEDULER->yield();

    SimpleDisk::issue_operation(DISK_OPERATION::WRITE, _block_no);
    done = false;
    issued = true;

    Console::puts("Wait for disk, yield\n");
    SYSTEM_SCHEDULER->yield();

    /* write data to port */
    int i;
    unsigned short tmpw;
    for (i = 0; i < 256; i++) {
        tmpw = _buf[2*i] | (_buf[2*i+1] << 8);
        Machine::outportw(0x1F0, tmpw);
    }

    wait_CPU = false;
    done = true;
    issued = false;
    Console::puts("WRITE DONE\n");
}
```


- After disk is ready, thread 3 finishes its read operation.

[illegible]